

Risk-averse Distributional Reinforcement Learning

A CVaR optimization approach

Silvestr Stanko¹

¹Department of Computer Science
Czech Technical University

Wednesday 30th May, 2018

Outline

Outline

Motivation



Figure: Robotics



Figure: Finance




Figure: AI safety

Ultimate goals of AI


General AI

- Learning from experience
- Learning *tabula rasa*
- Beyond purpose-specific AI
- Beyond human-level performance

Safe AI

- Avoiding catastrophic events
- Robust to environment changes or adversaries

Reinforcement Learning



Recent successes




Figure: Atari games

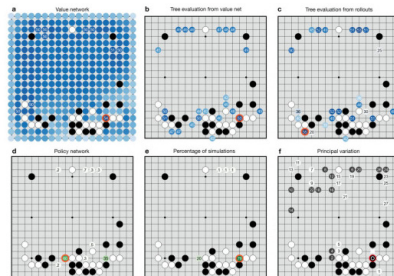


Figure: Go

Markov Decision Processes

Definition

An MDP is a 5-tuple $\mathcal{M} = (\mathcal{X}, \mathcal{A}, R, P, \gamma)$, where

- \mathcal{X} is the state space
- \mathcal{A} is the action space
- $R(x, a)$ is a random variable representing the reward generated by being in state x and selecting action a
- $P(\cdot|x, a)$ is the transition probability distribution
- $\gamma \in [0, 1)$ is a discount factor

Markov Decision Process - example

Reinforcement Learning - Goal

Definition

$Z^\pi(x_t)$ Is a random variable representing the discounted reward along a trajectory generated by the MDP by following the policy π , starting at state x_t .

$$Z^\pi(x_t) = \sum_{t=0}^{\infty} \gamma^t R(x_t, \pi(x_t))$$

Reinforcement Learning goals

Our goal is to find a globally optimal policy π^*

$$\pi^* = \arg \max_{\pi} \exp Z^\pi(x_0)$$

Potential problems

Problems

- Solutions must avoid catastrophic events and be **safe**
- RL is sample inefficient \rightarrow expensive training
- Solutions must be **robust** to small model changes

Solution

Instead of maximizing the expected reward, focus on other criteria that take into account the **risk** of the potential reward.

Risk

Definition

Risk is the potential of gaining or losing something of value.

Risk-averse: disinclined or reluctant to take risks

Risk-neutral: indifferent to or balanced with respect to risk.

Risk-seeking: inclined or eager to take risks

Example

Choose between receiving:

- 1 \$100 in 100% cases
- 2 \$200 in 50% cases and \$0 in 50% cases
- 3 \$10,000 in 1% cases and \$0 in 99% cases

Measuring Risk

Value-at-Risk (VaR)

- Easy to understand
- Historically the most used risk-measure
- Undesirable computational properties
- Does not differentiate between large and catastrophic losses

Definition

Let Z be a random variable representing reward, with cumulative distribution function $F(z) = \mathbb{P}(Z \leq z)$. The Value-at-Risk at confidence level $\alpha \in (0, 1)$ is the α -quantile of Z , i.e.

$$\text{VaR}_\alpha(Z) = F^{-1}(\alpha) = \inf \{z | \alpha \leq F(z)\}$$

Measuring Risk

Conditional Value-at-Risk (CVaR)


- Good computational properties
- Basel Committee on Banking Supervision: $\text{VaR} \rightarrow \text{CVaR}$
- Equivalent to robustness

Definition

The Conditional Value-at-Risk (CVaR) at confidence level $\alpha \in (0, 1)$ is defined as the expected reward of outcomes worse than the α -quantile (VaR_α):

$$\text{CVaR}_\alpha(Z) = \frac{1}{\alpha} \int_0^\alpha F_Z^{-1}(\beta) d\beta = \frac{1}{\alpha} \int_0^\alpha \text{VaR}_\beta(Z) d\beta$$

Value-at-Risk, Conditional Value-at-Risk



Conditional Value-at-Risk as an optimal point

Definition

$$\text{CVaR}_\alpha(Z) = \max_s \left\{ \frac{1}{\alpha} \mathbb{E} [(Z - s)^-] + s \right\}$$

where $(x)^- = \min(x, 0)$ and in the optimal point it holds that $s^* = \text{VaR}_\alpha(Z)$

$$\text{CVaR}_\alpha(Z) = \frac{1}{\alpha} \mathbb{E} [(Z - \text{VaR}_\alpha(Z))^-] + \text{VaR}_\alpha(Z)$$

it's dual is

$$\text{CVaR}_\alpha(Z) = \min_{\xi \in \mathcal{U}_{\text{CVaR}}(\alpha, p(\cdot))} \mathbb{E}_\xi[Z]$$

$$\mathcal{U}_{\text{CVaR}}(\alpha, p(\cdot)) = \left\{ \xi : \xi(z) \in \left[0, \frac{1}{\alpha}\right], \int \xi(z) p(z) dz = 1 \right\}$$

Risk-averse Reinforcement Learning - goals

Definition

$Z^\pi(x_t)$ is a random variable representing the discounted reward along a trajectory generated by the MDP by following the policy π , starting at state x_t .

$$Z^\pi(x_t) = \sum_{t=0}^{\infty} \gamma^t R(x_t, \pi(x_t))$$

Reinforcement Learning with CVaR

For a given α , our goal is to find a globally optimal policy π^*

$$\pi^* = \arg \max_{\pi} CVaR_{\alpha}(Z^\pi(x_0))$$

Risk-averse Reinforcement Learning - example

Figure: Greedy agent

Figure: Risk-averse agent

Outline

Value Iteration

Definition

Value function $V(x)$ represents the expected return when starting in state x and following the optimal policy π^* thereafter.

Value Iteration

Initialize $V_0(x)$ for each state (arbitrary value, e.g. 0).

Update each state:

$$V_{k+1}(x) = \max_a \left[R(x, a) + \gamma \sum_{x'} p(x'|x, a) V_k(x') \right]$$

Repeat.



The algorithm converges to the optimal policy π^* : $\lim_{k \rightarrow \infty} V_k(x) = V(x)$

Value Iteration with CVaR

Theorem (CVaR decomposition)

The conditional CVaR under policy π obeys the following decomposition:

$$CVaR_{\alpha}(Z^{\pi}(x, a)) = \min_{\xi \in \mathcal{U}_{CVaR}(\alpha, p(\cdot|x, a))} \sum_{x'} p(x'|x, a) \xi(x') CVaR_{\xi(x')\alpha}(Z^{\pi}(x'))$$



CVaR Value Iteration

Theorem (CVaR Value Iteration)

The following Bellman operator is a contraction:

$$\mathbf{T}C(x, y) = \max_a \left[R(x, a) + \gamma \min_{\xi} \sum_{x'} p(x'|x, a) \xi(x') C(x', y\xi(x')) \right]$$

The operator \mathbf{T} describes the following relationship:

$$\mathbf{T}CVaR_y(Z(x)) = \max_a \left[R(x, a) + \gamma CVaR_y(Z(x')) \right]$$

$$x' \sim p(\cdot|x, a)$$

Linear interpolation

Computing operator \mathbf{T} is intractable, as the state-space is continuous. A solution would be to approximate the operator with linear interpolation.

Theorem


The function $\alpha CVaR_\alpha$ is convex. The operator $\mathbf{T}_\mathcal{I}$ is a contraction.

$$\mathcal{I}_x[C](y) = y_i C(x, y_i) + \frac{y_{i+1} C(x, y_{i+1}) - y_i C(x, y_i)}{y_{i+1} - y_i} (y - y_i)$$

$$\mathbf{T}_\mathcal{I} C(x, y) = \max_a \left[R(x, a) + \gamma \min_\xi \sum_{x'} p(x'|x, a) \frac{\mathcal{I}_{x'}[C](y\xi(x'))}{y} \right]$$

This iteration can be formulated and solved as a linear program.

Linear interpolation



Original Contributions

① Faster CVaR Value Iteration

- Polynomial \rightarrow linear time.
- Formally proved for increasing, unbounded distributions.
- Experimentally verified for general distributions.

② CVaR Q-learning

- Sampling version of CVaR Value Iteration.
- Based on the distributional approach.
- Experimentally verified.

③ Distributional Policy improvement

- Proved monotonic improvement for distributional RL.
- Used as a heuristic for extracting π^* from CVaR Q-learning.

④ Deep CVaR Q-learning

- TD update \rightarrow loss function.
- Experimentally verified in a deep learning context.

CVaR VI computational complexity

Problem

- CVaR Value Iteration requires computing a Linear Program for each state and atom.
- LP computation is slow

Solution

CVaR Value Iteration with quantile representation.

αCVaR_α describes a quantile function

Lemma

Any discrete distribution has a piece-wise linear αCVaR_α function. Similarly, any a piece-wise linear αCVaR_α function can be seen as representing a certain discrete distribution.





$$\alpha\text{CVaR}_\alpha \Rightarrow \text{VaR}$$

$$\frac{\partial}{\partial \alpha} \alpha\text{CVaR}_\alpha(Z) = \frac{\partial}{\partial \alpha} \int_0^\alpha \text{VaR}_\beta(Z) d\beta = \text{VaR}_\alpha(Z)$$

$$\alpha\text{CVaR}_\alpha \Leftarrow \text{VaR}$$


$$\alpha\text{CVaR}_\alpha(Z) = \int_0^\alpha \text{VaR}_\beta(Z) d\beta$$

Next state CVaR computation




0.25

0.75




0.25

0.75



Next state CVaR computation



Linear-time Computation

Theorem

Solution to minimization problem

$$\min_{\xi \in \mathcal{U}_{CVaR}(\alpha, p(\cdot|x, a))} \sum_{x'} p(x'|x, a) \xi(x') CVaR_{\xi(x')\alpha} (Z^\pi(x'))$$





can be computed by setting

$$\xi(x') = \frac{F_{Z(x')}(F_{Z(x,a)}^{-1}(\alpha))}{\alpha}$$

The computational complexity is $O(n \cdot m)$ where n is the number of transition states and m is the number of atoms.

- Proved for increasing unbounded distributions

CVaR Value Iteration - Experiments



Outline

Q-learning

Problem

- In practice, we often don't have access to the transition probabilities $p(x'|x, a)$
- We need to learn through direct interaction with the environment.

Solution

Q-learning: Sampling version of Value Iteration.

Q-learning

Value Iteration

$$Q(x, a) \leftarrow \mathcal{T}Q(x, a)$$

$$Q(x, a) \leftarrow R(x, a) + \gamma \sum_{x'} p(x'|x, a) \max_{a'} Q(x', a')$$

Q-learning

$$Q(x, a) \leftarrow (1 - \beta)Q(x, a) + \beta \mathcal{T}Q(x, a)$$

$$Q(x, a) \leftarrow (1 - \beta)Q(x, a) + \beta \left[R(x, a) + \gamma \max_{a'} Q(x', a') \right]$$

CVaR estimation

Requirements:

- Store a single value
- Expected value of updates is CVaR

Recursive CVaR Estimation

$$V_{t+1} = V_t + \beta_t \left[1 - \frac{1}{\alpha} \mathbb{1}_{(V_t \geq r)} \right]$$
$$C_{t+1} = (1 - \beta_t) C_t + \beta_t \left[V_t + \frac{1}{\alpha} (r - V_t)^- \right]$$

CVaR Q-learning

Pseudocode

- 1 Sample a transition x, a, x', r
- 2 Create a target distribution \mathbf{d}
- 3 Sample from the target distribution
- 4 Update current estimates of VaR and CVaR towards the sample

Improvement

- ~~Sample from target distribution~~
- Update proportionally to the target distribution

CVaR Q-learning

CVaR Q-learning: Uniform case

```

1: input:  $x, a, x', r$ 
2: for each  $i$  do
3:    $C(x', y_i) = \max_{a'} C(x', a', y_i)$ 
4: end for
5:  $\mathbf{d} = \text{extractDistribution}(C(x', \cdot), \mathbf{y})$ 
6: for each  $i, j$  do
7:    $V(x, a, y_i) = V(x, a, y_i) + \beta \left[ 1 - \frac{1}{y_i} \mathbb{1}_{(V(x, a, y_i) \geq r + \gamma d_j)} \right]$ 
8:    $C(x, a, y_i) =$ 
       $(1 - \beta)C(x, a, y_i) + \beta \left[ V(x, a, y_i) + \frac{1}{y_i} (r + \gamma d_j - V(x, a, y_i))^- \right]$ 
9: end for

```

CVaR Q-learning

CVaR Q-learning: General case

```

1: input:  $x, a, x', r$ 
2: for each  $i$  do
3:    $C(x', y_i) = \max_{a'} C(x', a', y_i)$ 
4: end for
5:  $\mathbf{d} = \text{extractDistribution}(C(x', \cdot), \mathbf{y})$ 
6: for each  $i$  do
7:    $V(x, a, y_i) = V(x, a, y_i) + \beta \mathbb{E}_j \left[ 1 - \frac{1}{y_i} \mathbb{1}_{(V(x, a, y_i) \geq r + \gamma d_j)} \right]$ 
8:    $C(x, a, y_i) =$ 
      $(1 - \beta)C(x, a, y_i) + \beta \mathbb{E}_j \left[ V(x, a, y_i) + \frac{1}{y_i} (r + \gamma d_j - V(x, a, y_i))^- \right]$ 
9: end for

```

Optimal policy extraction

Standard RL: Optimal policy

$$\pi^*(x) = \arg \max_a Q(x, a)$$

CVaR VI: Optimal policy

$$\pi^*(x_0) = \arg \max_a C(x, a, \alpha)$$

$$\pi^*(x_1) = \arg \max_a C(x_1, a, \alpha \xi^*(x_0))$$

$$\vdots$$

$$\pi^*(x_t) = \arg \max_a C(x_t, a, y_{t-1} \xi^*(x_{t-1}))$$

Optimal policy extraction

Problem

- To compute y_t , we would need to know ξ^* .
- To compute ξ^* , we would need to know the probability of transitions $p(x'|x, a)$

Solution

- Use transition reward to compute the next-state y

Distributional Policy Improvement

CVaR primal definition:

$$\text{CVaR}_\alpha(Z) = \max_s \left\{ \frac{1}{\alpha} \mathbb{E} [(Z - s)^-] + s \right\}$$

Our goal can then be rewritten as

$$\max_{\pi} \text{CVaR}_\alpha(Z^\pi) = \max_{\pi} \max_s \frac{1}{\alpha} \mathbb{E} [(Z^\pi - s)^-] + s$$

Recall: The primal solution is equivalent to $\text{VaR}_\alpha(Z)$

Idea: If we knew the value $s^* = \text{VaR}_\alpha(Z)$ in advance, we could simplify the problem to maximize only

$$\max_{\pi} \text{CVaR}_\alpha(Z^\pi) = \max_{\pi} \frac{1}{\alpha} \mathbb{E} [(Z^\pi - s^*)^-] + s^*$$

Distributional Policy Improvement

Policy Improvement

- 1 Maximize $\frac{1}{\alpha} \mathbb{E} [(Z^\pi(x_0) - s)^-] + s$ w.r.t. s while keeping π fixed.
- 2 Maximize $\frac{1}{\alpha} \mathbb{E} [(Z^\pi(x_0) - s)^-] + s$ w.r.t. π while keeping s fixed.
- 3 Recompute $\text{CVaR}_\alpha(Z^{\pi^*})$ where π^* is the new policy.

VaR-based Policy Improvement

Theorem

Let π be a fixed policy, $\alpha \in (0, 1]$. By following policy π' from the following algorithm, we will improve $CVaR_\alpha(Z)$ in expectation:

$$CVaR_\alpha(Z^\pi) \leq CVaR_\alpha(Z^{\pi'})$$

VaR-based Policy Improvement

$a = \arg \max_a CVaR_\alpha(Z(x_0, a))$

$s = VaR_\alpha(Z(x_0, a))$

Take action a , observe x, r

while x is not terminal **do**

$$s = \frac{s - r}{\gamma}$$

$a = \arg \max_a \mathbb{E}[(Z(x, a) - s)^-]$

Take action a , observe x, r

end while


VaR-based heuristic

Problem





When using quantile discretization, we don't have access to the exact VaR .

Solution

Use linear interpolation as a heuristic.



CVaR Q-learning - Experiments



Outline

Approximate Q-learning

Problem

Q-learning is intractable for large state spaces.

Solution

Use approximate Q-learning.

- Formulate CVaR Q-learning update as a minimizing argument
- Use methods of convex optimization to find the optimal point

TD update \rightarrow loss function

Standard RL

$$Q(x, a) = (1 - \beta)Q(x, a) + \beta \mathcal{T}Q(x, a)$$

$$\downarrow$$

$$\min_{\theta} \mathbb{E} \left[(Q_{\theta}(x, a) - \mathcal{T}Q(x, a))^2 \right]$$

CVaR RL


$$V(x, a, y_i) = V(x, a, y_i) + \beta \mathbb{E}_j \left[1 - \frac{1}{y_i} \mathbb{1}_{(V(x, a, y_i) \geq \mathcal{T}d_j)} \right]$$

$$C(x, a, y_i) = (1 - \beta)C(x, a, y_i) + \beta \mathbb{E}_j \left[V(x, a, y_i) + \frac{1}{y_i} (r + \gamma d_j - V(x, a, y_i))^- \right]$$

$$\downarrow$$

$$???$$

Quantile loss



TD update \rightarrow loss function

VaR loss

$$\mathcal{L}_{\text{VaR}} = \sum_{i=1}^N \mathbb{E}_j \left[(r + \gamma d_j - V_i(x, a))(y_j - \mathbb{1}_{(V_i(x, a) \geq r + \gamma d_j)}) \right]$$

CVaR loss


$$\mathcal{L}_{\text{CVaR}} = \sum_{i=1}^N \mathbb{E}_j \left[\left(V_i(x, a) + \frac{1}{y_i} (r + \gamma d_j - V_i(x, a))^- - C_i(x, a) \right)^2 \right]$$

$$\mathcal{L} = \mathbb{E} [\mathcal{L}_{\text{VaR}} + \mathcal{L}_{\text{CVaR}}]$$

Deep CVaR Q-learning

- Model: Convolutional Neural Network
 - 1 Input: $84 \times 84 \times 4$
 - 2 Convolution: $8 \times 8 \times 32$ (stride 4)
 - 3 Convolution: $4 \times 4 \times 64$ (stride 2)
 - 4 Convolution: $3 \times 3 \times 32$ (stride 1)
 - 5 Fully connected: 256 hidden units
 - 6 Output: $|\mathcal{A}| \times 100$
- Replay Memory
- Target network C'
- Optimizer: Adam (Stochastic Gradient Descent)

Deep CVaR Q-learning - Experiments



- ① Video: $\alpha = 1$
- ② Video: $\alpha = 0.3$

Summary

① Faster CVaR Value Iteration

- Polynomial \rightarrow linear time.
- Formally proved for increasing, unbounded distributions.
- Experimentally verified for general distributions.

② CVaR Q-learning

- Sampling version of CVaR Value Iteration.
- Based on the distributional approach.
- Experimentally verified.

③ Distributional Policy improvement

- Proved monotonic improvement for distributional RL.
- Used as a heuristic for extracting π^* from CVaR Q-learning.

④ Deep CVaR Q-learning

- TD update \rightarrow loss function.
- Experimentally verified in a deep learning context.

Future work

Theory

- CVaR Value Iteration - General equivalence proof
- Q-learning - convergence proof

Practice

- Larger state spaces
- Practical problems (e.g. finance)